

CS344 Assignment-2

Group-3, CSE

August 18, 2025

Team Members

- 230101024 (Baswa Dhanush Sai)
- 230101114 (Yeredla Koushik Reddy)
- 230101092 (Seella Mythresh Babu)
- 230101120 (Taes Padhiary)

Please find the Google Drive folder containing all the edited files here:

drive.google.com/drive/folders/1U15WB5L0lYgSJm6dPksa7rLtmEL7pat2?usp=sharing

Contents

1 Task 2.1: Weighted Round Robin Scheduler	1
1.1 Problem Description	1
1.2 Implementation in xv6	2
1.2.1 kernel and user changes	2
1.3 Testing with task2.1Demo.c	3
2 Task 2.2: Custom Scheduling Algorithm	4
2.1 Problem Description	4
2.2 Implementation in xv6	4
2.2.1 Design and Tradeoffs: Lottery Scheduling	5
2.3 Testing with task2.2Demo.c	6

[1] Task 2.1: Weighted Round Robin Scheduler

1.1. Problem Description

The default xv6 scheduler treats all processes equally. In this task, modify it to allow processes with higher priority to get more CPU time, while ensuring lower priority processes are not completely starved.

1. Add new system calls to allow setting and getting process priorities:
 - `set_priority(n)`: sets priority of the current process to a positive integer `n`.
 - `get_priority()`: returns priority of the calling process. MAX value is 1000.
2. Modify the xv6 scheduler to implement a Weighted Round Robin (WRR) algorithm using the priority values as weights:
 - A process with higher priority should get proportionally more CPU time than one with lower priority.
 - However, no process should be starved; all runnable processes must get some CPU time.
 - Set a default priority for all new processes (e.g., 10), unless explicitly changed.

1.2. Implementation in xv6

1. **Per-process add-ons:**
 - `priority` (default: 10, max: 1000) — higher means longer quantum.
 - a per-run counter (`ticks`) that counts how many timer ticks the process has used in its current quantum. Initialized to 0.
2. **Initialization:** New processes start with `priority=DEFAULT_PRIORITY` and begin a fresh quantum whose target length equals `priority`.
3. **Preemption** (main edit): In both `usertrap()` and `kerneltrap()`, on each timer interrupt we:
 - (a) increment the running process' `p->ticks`
 - (b) if `p->ticks >= p->priority`, reset `p->ticks=0` and `yield()` (end of the time quantum)
 - (c) otherwise, **do not** yield (the process continues running across ticks).
4. **System calls:** `setpriority(n)`, `getpriority()` allow user-space to set/read the weight. When setting a new priority we immediately reset the quantum.

1.2.1. kernel and user changes

- We define `MAX_PRIORITY = 1000` and `DEFAULT_PRIORITY = 10` in `kernel/defs.h`.
- In `kernel/proc.h`, we extend `struct proc` with `priority` (the WRR weight) and `ticks` (the number of ticks this process has consumed in the *current* quantum).
- In `kernel/proc.c`, set `p->priority=DEFAULT_PRIORITY` and `p->ticks=0` in the `allocproc()` and set `p->ticks=0` in `fork()` so it begins with a new quantum.

- (Core idea) • In `kernel/trap.c`: xv6 originally *yields* on *every* timer tick — we change that to *yield only when the current process has used priority ticks* in both `usertrap()` and `kerneltrap()`.
- Add the system calls `int setpriority(int n)` and `int getpriority(void)` as numbers in `kernel/syscall.h`, tables in `kernel/syscall.c`, implementations in `kernel/sysproc.c` and prototypes in `user/user.h` and `user/user.pl`.

1.3. Testing with `task2.1Demo.c`

The following code (output in Figure 1) checks whether Weighted Round Robin is enforced. ‘A’ has been given priority 500 and hence can run upto 500 ticks uninterrupted. Then ‘B’ will run till the time quantum expires and will be preempted by ‘C’ which will run till 100 ticks/remaining time quanta uninterrupted if no higher priority process is present.

Listing 1. `user/task2.1Demo.c`

```
#include "kernel/types.h"
#include "user/user.h"

int priority_fork(int priority, char label) {
    int p = fork();
    if (p == 0) {
        setpriority(priority);

        for (long i = 0; i < 1000000000; i++) {
            if (i % 10000000 == 0) write(1, &label, 1);
        }

        fprintf(1, "\nProcess %d finished\n", priority);
        exit(0);
    }
    return p;
}

int main(void) {

    priority_fork(500, 'A');
    priority_fork(1, 'B');
    priority_fork(100, 'C');

    wait(0);
    wait(0);
    wait(0);

    exit(0);
}
```

```

Ubuntu-24.04
BOL TABLE/d; s/ .* / /; / ^$/d' > user/task2.1Demo.sym
mkfs/mkfs fs.img README user/_cat user/_echo user/_forktest us
er/_grep user/_init user/_kill user/_ln user/_ls user/_mkdir u
ser/_rm user/_sh user/_stressfs user/_usertests user/_grind us
er/_wc user/_zombie user/_task2.1Demo
nmeta 46 (boot, super, log blocks 30, inode blocks 13, bitmap
blocks 1) blocks 1954 total 2000
ballocc: first 823 blocks have been allocated
ballocc: write bitmap block at sector 45
qemu-system-riscv64 -machine virt -bios none -kernel kernel/ke
rnel -m 128M -smp 1 -nographic -global virtio-mmio.force-legac
y=false -drive file=fs.img,if=none,format=raw,id=x0 -device vi
rtio-blk-device,drive=x0,bus=virtio-mmio-bus.0

xv6 kernel is booting

init: starting sh
$ task2.1Demo
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Process 500 finished
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB
Process 100 finished
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB
Process 1 finished
$

```

Figure 1. ‘A’ [500] runs till either 500 ticks or till it is completed (here, the latter happens). ‘B’ [1] runs till the end of time slice and then gets pre-empted by ‘C’ [100] which runs till it is finished. Then, process ‘B’ finishes uninterrupted. Hence, weighted Round Robin is enforced here.

[2] Task 2.2: Custom Scheduling Algorithm

2.1. Problem Description

Move beyond round-robin and weighted round-robin scheduling. Analyse and implement a scheduling algorithm that optimises one or more of the following goals:

- Minimise turnaround or response time
- Avoid starvation
- Improve CPU utilisation
- Ensure fairness
- Reduce waiting time

2.2. Implementation in xv6

We will implement LOTTERY SCHEDULING. The repository at github.com/avaiyang/xv6-lottery-scheduling provided information on the steps to be followed to implement the lottery scheduling algorithm.

2.2.1. Design and Tradeoffs: Lottery Scheduling

1. **Per-process state:** each process `p` stores an integer `tickets` (≥ 1). New processes start with `tickets` = 10 by default. Children inherit the parent's tickets.
2. **Syscall:** `settickets(int n)` sets the caller's `tickets` to `n` if $n \geq 1$.
3. **Main scheduling loop (runs forever):**
 - Set `total` to 0. Scan the process table:
For every `RUNNABLE` process `p`, let $t = \max(1, p \rightarrow \text{tickets})$ and add `t` to `total`.
 - If `total == 0`, no process is runnable: execute `wfi` and restart the loop.
 - Draw a winning number `w` by computing `w = rand32() % total`.
 - Set `prefix` = 0. Scan the process table again:
 - For each `RUNNABLE` process `p`, let $t = \max(1, p \rightarrow \text{tickets})$.
 - If $w < \text{prefix} + t$, this process owns the winning ticket:
 - * Set `p->state` = `RUNNING`, set `c->proc` = `p`, and perform `swtch` to `p`.
 - * When `swtch` returns, clear `c->proc` and break out to start a new scheduling decision.
 - Otherwise, increase `prefix` by `t` and continue scanning.
4. **PRNG used by the scheduler (`rand32()`):**
 - Maintain a 64-bit nonzero global (or per-CPU) state `state`.
 - On each call: `state = state * 6364136223846793005ULL + 1ULL` (wraps mod 2^{64}). Return the upper 31 bits using `(uint)(state >> 33)`. This is the winner lottery value.
5. **Correctness:**
 - Over time the CPU share of process i converges to almost $t_i / \sum t_i$.
 - Any process with at least one ticket has non-zero probability of being chosen at each decision.
 - User space controls weights via `settickets` (e.g., interactive jobs should ideally get more tickets).
6. **Drawbacks:**
 - Random selection yields bursty service with the proportions stabilizing only over longer windows.
 - Without caps, a process can inflate its tickets and dominate which can lead to a potential abuse of the scheduler, as the user can set the tickets.
 - $O(NPROC)$ per decision is not scalable.
 - Even with high tickets, a process can occasionally experience unlucky gaps; not suitable where strict minimum service guarantees are required.

2.3. Testing with task2.2Demo.c

Listing 2. user/task2.2Demo.c

```
#include "kernel/types.h"
#include "kernel/stat.h"
#include "user/user.h"

// a CPU-bound function
void cpu_work(long iter) {
    for (long i = 0; i < iter; i++) {
        // A volatile variable prevents the compiler from optimizing the
        // loop away.
        volatile int x = 1 + 2*3 - 4 / 1 + 91 - 35 * 1729 + 8;
        (void)x; // Suppress "unused variable" warning.
    }
}

// function to test lottery scheduling
int fork_lottery_child(const char *label, int tickets, int dur) {
    int p = fork();
    if (p==0) {
        if (settickets(tickets) < 0) printf("%s: settickets failed\n",
            label);

        uint64 start = uptime(), cnt = 0;
        while(uptime() < start + dur){
            cpu_work(5000);
            cnt++;
        }

        printf("%s tickets=%d cnt=%d\n", label, tickets, (int)cnt); // "
            cnt" will be proportional to the number of "tickets"
        exit(0);
    }
    return p;
}

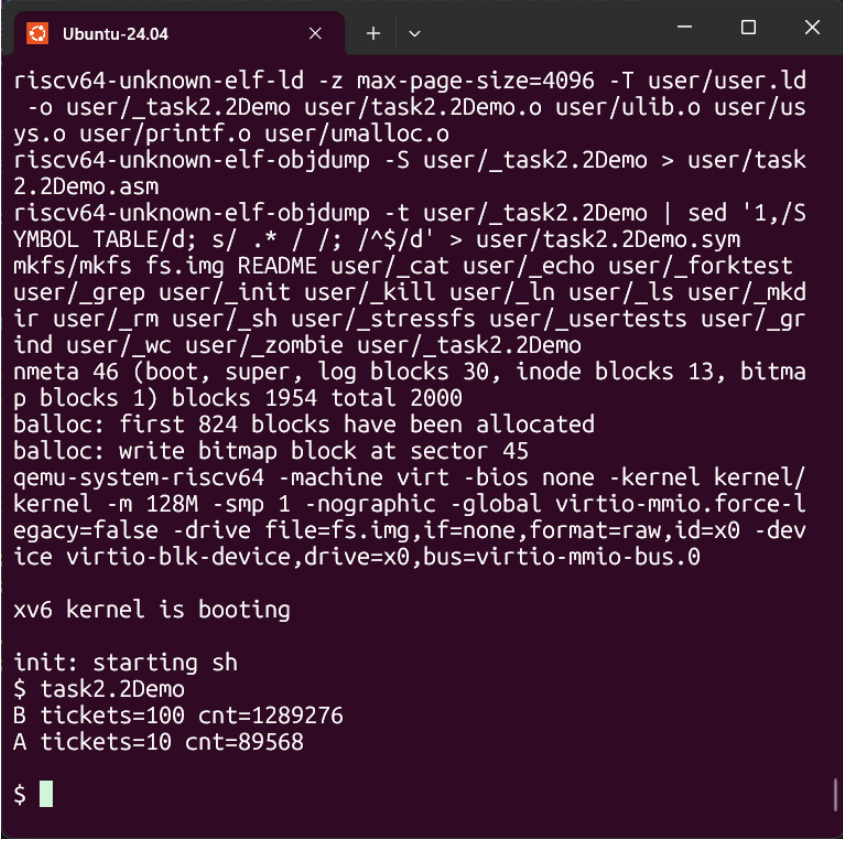
int
main(void)
{
    int t1 = 10;
    int t2 = 100;
    int dur = 300;

    fork_lottery_child("A", t1, dur);
    fork_lottery_child("B", t2, dur);

    wait(0);
    wait(0);

    write(1, "\n", 1);
    exit(0);
}
```

The output in Figure 2 provides the expected outcome of the Lottery scheduling involving two processes.



```
Ubuntu-24.04
riscv64-unknown-elf-ld -z max-page-size=4096 -T user/user.ld
-o user/_task2.2Demo user/task2.2Demo.o user/ulib.o user/us
ys.o user/printf.o user/umalloc.o
riscv64-unknown-elf-objdump -S user/_task2.2Demo > user/task
2.2Demo.asm
riscv64-unknown-elf-objdump -t user/_task2.2Demo | sed '1,/S
YMBOL TABLE/d; s/ .* / /; /^$/d' > user/task2.2Demo.sym
mkfs/mkfs fs.img README user/_cat user/_echo user/_forktest
user/_grep user/_init user/_kill user/_ln user/_ls user/_mkd
ir user/_rm user/_sh user/_stressfs user/_usertests user/_gr
ind user/_wc user/_zombie user/_task2.2Demo
nmeta 46 (boot, super, log blocks 30, inode blocks 13, bitma
p blocks 1) blocks 1954 total 2000
ballocc: first 824 blocks have been allocated
ballocc: write bitmap block at sector 45
qemu-system-riscv64 -machine virt -bios none -kernel kernel/
kernel -m 128M -smp 1 -nographic -global virtio-mmio.force-l
egacy=false -drive file=fs.img,if=none,format=raw,id=x0 -dev
ice virtio-blk-device,drive=x0,bus=virtio-mmio-bus.0

xv6 kernel is booting

init: starting sh
$ task2.2Demo
B tickets=100 cnt=1289276
A tickets=10 cnt=89568

$
```

Figure 2. Fork labelled 'A' with 10 tickets ran around 89568 times, while fork labelled 'B' ran 1289276 ($> 11 \cdot 89568$) times, thus proving the expected lottery scheduling