

# CS344 Assignment 5

## Group-3, CSE

October 8, 2025

## Team Members

- **230101024** (Baswa Dhanush Sai)
- **230101092** (Seella Mythresh Babu)
- **230101114** (Yeredla Koushik Reddy)
- **230101120** (Taes Padhihary)

Please find the Google Drive folder containing all the edited files here:

[drive.google.com/drive/folders/106CkVtq2gCoXgV4hyZ2kui2GwBto0w\\_0?usp=sharing](https://drive.google.com/drive/folders/106CkVtq2gCoXgV4hyZ2kui2GwBto0w_0?usp=sharing)

## Contents

<b>1 Task 5.1: Doubly-Indirect Block Support</b>	<b>1</b>
1.1 Implementation . . . . .	1
1.2 Testing Doubly-Indirect Block Support . . . . .	4
<b>2 Task 5.2: Symbolic Links (Soft Links)</b>	<b>4</b>
2.1 Implementation . . . . .	4
2.2 Testing Symbolic links . . . . .	6

---

## 1. Task 5.1: Doubly-Indirect Block Support

**Reference used:** [xiayingp.gitbook.io/build\\_a\\_os/labs/lab-8-file-system-large-files](https://xiayingp.gitbook.io/build_a_os/labs/lab-8-file-system-large-files)

**Synopsis of the problem** Figure 1 shows the format of storing files in xv6 that is currently used. It supports  $12 + 256 = 268$  blocks. This limits the total size of the files to around 268KiB. We need to increase this size.

### 1.1. Implementation

We need to support a “doubly-indirect” block in each inode, containing 256 addresses of singly indirect blocks, each of which can further contain upto 256 blocks. The new plan is shown in Figure 2. This increases the size to  $11 + 256 + 256 \cdot 256 = 65803$  blocks.




Figure 8.3: The representation of a file on disk.

Figure 1. Old file representation (source: [xv6/book-riscv-rev4.pdf](http://xv6/book-riscv-rev4.pdf))

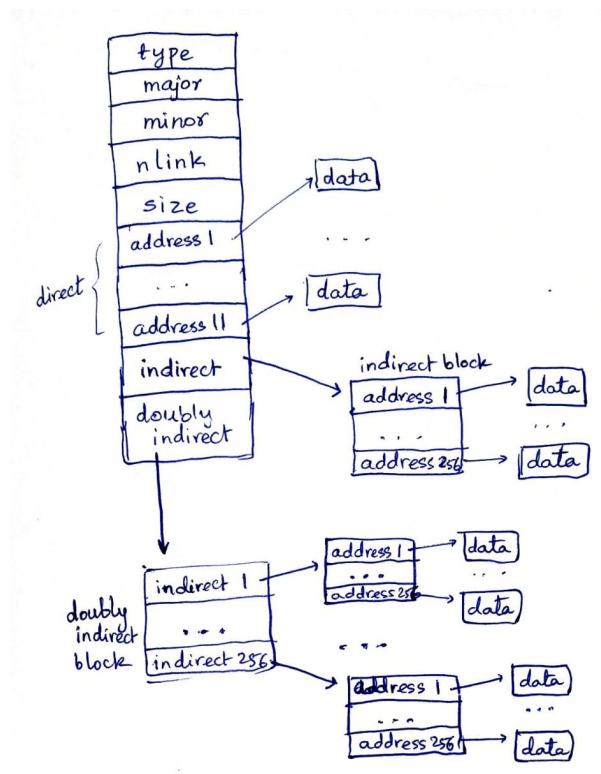


Figure 2. Implemented file representation

## Mechanism

1. Update the **FSSIZE** macro in `kernel/param.h` to a large number like **200000** instead of 2000.
2. In `kernel/fs.h`, we update the value of **NDIRECT** to 11 (from the initial value of 12, because we should not change the disk `addrs` size), add the macro  
**NDOUBLY\_INDIRECT** = **NINDIRECT** \* **NINDIRECT** and update  
**MAXFILE** = **NDIRECT** + **NINDIRECT** + **NDOUBLY\_INDIRECT**.
3. In `kernel/file.h`, change the size of `addrs[]` array of structure `inode` from **NDIRECT**+1 to **NDIRECT**+2 (`addrs[NDIRECT]` will have the indirect block and `addrs[NDIRECT+1]` will have the doubly-indirect block of Figure 2). This is the *in-memory* copy of the inode structure. Make the same changes to the *on-disk* copy by updating the size of `addrs[]` array to **NDIRECT**+2 in `kernel/fs.h`.
4. Finally, to implement doubly-ended indirection of files, we modify the functions `bmap()` and `itrunc()` of `kernel/fs.c` as shown in Listing 2.

**Listing 1.** changes in macros and structs

```
/* kernel/param.h */
#define FSSIZE 200000 // initially 2000, but
                      now 200000

/* kernel/file.h */
// ...
struct inode {
    // ...
    uint addrs[NDIRECT+2];
};

/* kernel/fs.h */
#define NDIRECT 11
#define NINDIRECT (BSIZE / sizeof(uint))
#define NDOUBLY_INDIRECT (NINDIRECT *
                          NINDIRECT)
#define MAXFILE (NDIRECT + NINDIRECT
                + NDOUBLY_INDIRECT)

struct dinode {
    // ...
    uint addrs[NDIRECT+2];
};
```

**Listing 2.** changes to `kernel/fs.c` for doubly-indirect block

```
/* kernel/fs.c */
static int bmap(struct inode *ip, uint bn){
    // ...
    if (bn < NDIRECT) { // calculate ... }
    // Task 5.1: handle doubly indirect blocks
    bn -= NDIRECT;
    if (bn < NDOUBLY_INDIRECT) {
        if ((addr = ip->addrs[NDIRECT+1]) == 0){
            addr = balloc(ip->dev);
            if (addr == 0) return 0;
            ip->addrs[NDIRECT+1] = addr;
        }
```

```
    }
    bp = bread(ip->dev, ip->addrs[NDIRECT
        +1]);
    a = (uint*)bp->data;
    // now, go to the second layer
    uint outer_index = bn / NINDIRECT;
    if ((addr = a[outer_index]) == 0){
        addr = balloc(ip->dev);
        if (addr == 0){ brelse(bp); return 0; }
        a[outer_index] = addr; log_write(bp);
    }
    brelse(bp);

    // the final page from disk
    uint inner_index = bn % NINDIRECT;
    bp = bread(ip->dev, addr);
    a = (uint*)bp->data;
    if ((addr = a[inner_index]) == 0){
        addr = balloc(ip->dev);
        if (addr == 0){ brelse(bp); return 0; }
        a[inner_index] = addr; log_write(bp);
    }
    brelse(bp); return addr;
}
}
```

```
void itrunc(struct inode *ip){
    int i, j, k;
    struct buf *bp, *bp2;
    uint *a, *a2;
    // ...
    // Task 5.1: free the doubly indirect blocks
    if (ip->addrs[NDIRECT + 1]){
        bp = bread(ip->dev, ip->addrs[NDIRECT
            +1]);
        a = (uint*)bp->data;
        for(j = 0; j < NINDIRECT; j++){
            if (a[j]){
                bp2 = bread(ip->dev, a[j]);
```

```

a2 = (uint*)bp2->data;
for(k = 0; k < NINDIRECT; k++)
    if (a2[k]) bfree(ip->dev, a2[k]);
    brelse(bp2); bfree(ip->dev, a[j]);
}
}

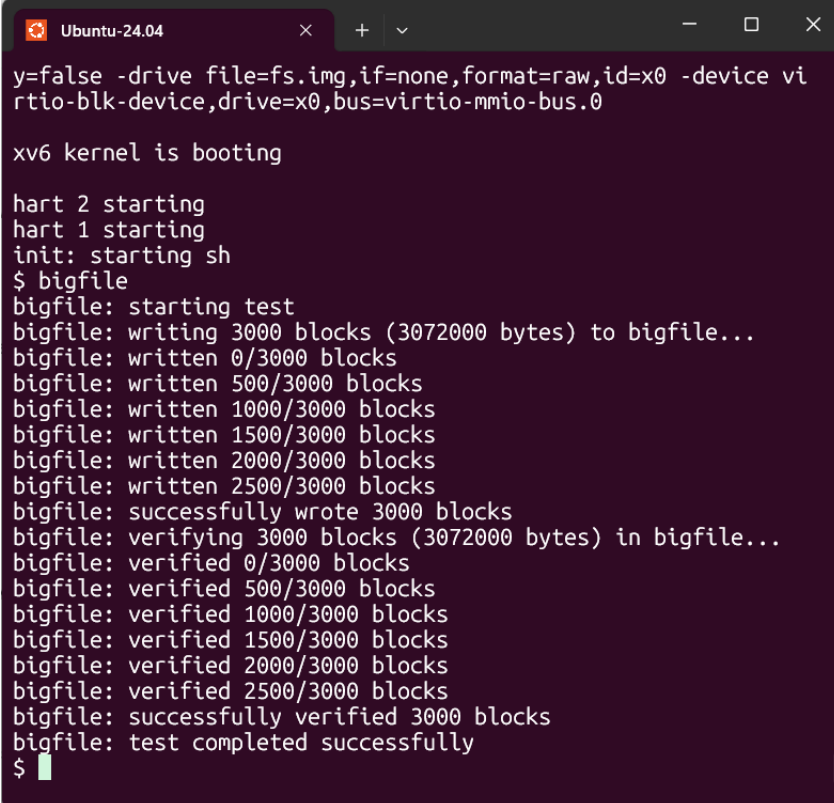
```

```

brelse(bp);
bfree(ip->dev, ip->addrs[NINDIRECT+1]);
ip->addrs[NINDIRECT+1] = 0;
}
}

```

## 1.2. Testing Doubly-Indirect Block Support



```

Ubuntu-24.04
y=false -drive file=fs.img,if=none,format=raw,id=x0 -device v
rtio-blk-device,drive=x0,bus=virtio-mmio-bus.0

xv6 kernel is booting

hart 2 starting
hart 1 starting
init: starting sh
$ bigfile
bigfile: starting test
bigfile: writing 3000 blocks (3072000 bytes) to bigfile...
bigfile: written 0/3000 blocks
bigfile: written 500/3000 blocks
bigfile: written 1000/3000 blocks
bigfile: written 1500/3000 blocks
bigfile: written 2000/3000 blocks
bigfile: written 2500/3000 blocks
bigfile: successfully wrote 3000 blocks
bigfile: verifying 3000 blocks (3072000 bytes) in bigfile...
bigfile: verified 0/3000 blocks
bigfile: verified 500/3000 blocks
bigfile: verified 1000/3000 blocks
bigfile: verified 1500/3000 blocks
bigfile: verified 2000/3000 blocks
bigfile: verified 2500/3000 blocks
bigfile: successfully verified 3000 blocks
bigfile: test completed successfully
$

```

**Figure 3.** The test file [user/bigfile.c](#) allocates 3000 blocks (file size), fills them with distinct values and then checks the values at those 3000 blocks. The figure shows successful completion of the test.

## 2. Task 5.2: Symbolic Links (Soft Links)

**Reference used:** [xiayingp.gitbook.io/build\\_a\\_os/labs/lab-8-file-system-symbolic-links](https://xiayingp.gitbook.io/build_a_os/labs/lab-8-file-system-symbolic-links)

**Synopsis of the problem** Symbolic links refer to a linked file by pathname; when a symbolic link is opened, the kernel follows the link to the referred file. Symbolic links resembles hard links, but hard links are restricted to pointing to file on the same disk, while symbolic links can cross disk devices.

### 2.1. Implementation

1. Add `T_SYMLINK` in [kernel/stat.h](#) to denote another type of inode.

2. Define `O_NOFOLLOW` in `kernel/fcntl.h`.
3. Define the system call number in `kernel/syscall.h` for `sys_symlink()` and add it to the `syscalls[]` table in `kernel/syscall.c`.
4. Do modifications to `sys_open()` and `sys_symlink()` in `kernel/sysfile.c` as shown in Listing 4.
  - `sys_symlink()`: Implements the `symlink` system call to create a symbolic link named `path` that points to `target`. Receives both arguments; calls `create(path, T_SYMLINK, 0, 0)` to allocate an inode of type `T_SYMLINK` and stores the `target` inside the inode's data blocks using `writei()`. It unlocks and releases the inode (`iunlockput(ip)`) and finishes the transaction (`end_op()`).
  - Modified `create()` to include inode of type `T_SYMLINK`.
  - `sys_open()`: Retrieve the pathname and its flags. If the inode type is `T_SYMLINK`: (1) if `O_NOFOLLOW` is specified, open the symbolic link file itself instead of the target; and (2) if `O_NOFOLLOW` is not specified, then recursively resolve the path (till the depth < 10) using `namei()` and `readi()` until either a non-link file is found or the depth  $\geq 10$ . Continue normal file descriptor allocation and setup once the correct inode is obtained.
5. Add the prototype of the syscall `symlink` in `user/user.h` and add it to `user/usys.pl`.

**Listing 3.** changes to variables

```
/* kernel/stat.h */
// ...
#define T_SYMLINK 4

/* kernel/fcntl.h */
// ...
#define O_NOFOLLOW 0x800

/* kernel/syscall.h */
// ...
#define SYS_symlink 22

/* kernel/syscall.c */
// ...
extern uint64 sys_symlink(void);

static uint64 (*syscalls[])(void) = {
    // ...
    [SYS_symlink] sys_symlink,
}

/* user/user.h */
// ...
int symlink(const char *target, const char *
            path);
```

**Listing 4.** changes to `sys_open()` and creating `sys_symlink()` in `kernel/sysfile.c`

```
static struct inode* create(char *path, short
                           type, short major, short minor){
    // ...
    if((ip = dirlookup(dp, name, 0)) != 0){
        iunlockput(dp);
        ilock(ip);
```

```
if((type == T_SYMLINK) || (type == T_FILE
    && (ip->type == T_FILE || ip->type ==
    T_DEVICE)))) // add T_SYMLINK
    recognition to inode -- Task 5.2
    return ip;
iunlockput(ip);
return 0;
}
// ...
}

// Task 5.2 --- syscall: sys_symlink
uint64 sys_symlink(void){
    char target[MAXPATH], path[MAXPATH];
    if (argstr(0, target, MAXPATH) < 0 || argstr(1,
        path, MAXPATH) < 0) return -1;

    begin_op();
    struct inode *ip = create(path, T_SYMLINK,
        0, 0);
    if (ip == 0){ end_op(); return -1; }

    int len = strlen(target);
    if(writei(ip, 0, (uint64)target, 0, len + 1) != (len
        + 1)){
        ip->nlink = 0; iupdate(ip); iunlockput(ip);
        end_op(); return -1;
    }

    iunlockput(ip);
    end_op();
    return 0;
}

uint64
sys_open(void)
```

```

{
//...
if(omode & O_CREATE){
    ip = create(path, T_FILE, 0, 0);
    if(ip == 0){ end_op(); return -1; }
} else {
    if((ip = namei(path)) == 0){ end_op(); return
        -1; }
    ilock(ip);
    // Task 5.2
    if ((ip->type == T_SYMLINK) && !(omode
        & O_NOFOLLOW)){
        int count = 0;
        while (ip->type == T_SYMLINK && count
            < 10){
            int len = 0;
            if (readi(ip, 0, (uint64)path, 0, len + 1) != (
                len+1)){ iunlockput(ip); end_op(); return
                -1; }

```

```

        iunlockput(ip);

        if((ip = namei(path)) == 0){ end_op();
            return -1; }
        ilock(ip);
        count++;
    }
    if (count >= 10) { // recursive depth till 10
        printf("cycle for symlink, error\n");
        iunlockput(ip);
        end_op();
        return -1;
    }
}
//...
}
//...

```

## 2.2. Testing Symbolic links

```

Ubuntu-24.04
riscv64-unknown-elf-objdump -t user/_symlinktest | sed '1,/SYM
BOL TABLE/d; s/ .* / /; / ^$/d' > user/symlinktest.sym
mkfs/mkfs fs.img README user/_cat user/_echo user/_forktest us
er/_grep user/_init user/_kill user/_ln user/_ls user/_mkdir u
ser/_rm user/_sh user/_stressfs user/_usertests user/_grind us
er/_wc user/_zombie user/_logstress user/_forphan user/_dorpha
n user/_symlinktest
nmeta 47 (boot, super, log blocks 31, inode blocks 13, bitmap
blocks 1) blocks 1953 total 2000
ballocc: first 959 blocks have been allocated
ballocc: write bitmap block at sector 46
qemu-system-riscv64 -machine virt -bios none -kernel kernel/ke
rnel -m 128M -smp 3 -nographic -global virtio-mmio.force-legac
y=false -drive file=fs.img,if=none,format=raw,id=x0 -device vi
rtio-blk-device,drive=x0,bus=virtio-mmio-bus.0

xv6 kernel is booting

hart 1 starting
hart 2 starting
init: starting sh
$ symlinktest
symlink follow OK
nofollow target='t.txt'
cycle for symlink, error
expected loop detection; successful test!
symlink tests done
$

```

**Figure 4.** Testing the implementation using `user/symlinktest.c`: We create a base file (`t.txt`), link it symbolically (`l.txt`), verify content access through the link, test the `O_NOFOLLOW` flag to read the target path itself, and finally check loop detection with recursive links ( $l_A \rightarrow l_B \rightarrow l_A$ ). The terminal output confirms all stages